

Data 188: Introduction to Deep Learning

Fully-connected networks, optimization, initialization

Speaker: Eric Kim
Lecture 06 (Week 03)
2026-02-05, Spring 2026. UC Berkeley.

Announcements

- HW0 due tonight!
- HW1 is released
 - Warning: this homework is substantially more work than HW0. Start early!



<https://phdcomics.com/comics/archive.php?comicid=820>



How to succeed in this class

- Step 1: Attend ALL course activities. **Don't fall behind!**
 - An easy way to ensure success: **regularly attend lecture and discussion**. Everything else will follow naturally.
 - It'll be very difficult to succeed in this class without doing at least this. "Why play on hard mode?"
- Step 2: Actively study material
 - Ex: "passive" learning would be simply watching lecture and moving on.
 - "Active" learning: revisit lecture slides and discussion exercises. Recreate any tricky derivations, create pictures/visualizations to reinforce concepts, ask questions on Ed, go to office hours, Google things...
 - **Make course concepts "yours".** "Can you explain ConceptX to someone not taking the course?"
- Step 3: Reinforce concepts by doing homework assignments
 - In this course, homework assignments are sort of the "first class citizens". In a sense, lectures and discussions exist to support you for the homework.
 - These homework assignments are a great way to get a deep, hands-on understanding of the course concepts
 - Homework assignments may reveal holes in your knowledge. This is natural! Forge on ahead, it'll be worth it!

Some good advice that you may enjoy: "[Doing well in your courses: a guide by Andrej Karpathy](#)".

Outline

Fully connected networks

Optimization

Initialization

Outline

Fully connected networks

Optimization

Initialization

Linear layer

In our MNIST digit classifier (one, two-layer), the primary operation is a matrix multiply (eg ops.MatMul):

One layer classifier

$$h(X) = X\theta$$

n: input dimensionality
k: number of output classes
d: hidden dimension size


X has shape=[batchsize, n], θ has shape=[n, k]

Two layer classifier

$$h(X) = \sigma(XW_1)W_2$$

W_1 has shape=[batchsize, d], W_2 has shape=[d, k]

In deep learning, it's common to augment MatMul with a learnable **bias** term. This is commonly called a "Linear" (or FullyConnected) layer:



$$\text{linear}(X) = XW + b$$

X shape=[batchsize, d_in]
W shape=[n, d_out]
b shape=[1, d_out]

Linear layer: illustrated

$\text{Linear}(X)$ takes an input with dimensionality `d_in` and transforms it to dimensionality `d_out`

$$W = \begin{bmatrix} 2 & 2 & 0 & 3 \\ 0 & 1 & 1 & 5 \\ 1 & 4 & 2 & 0 \end{bmatrix}$$
$$b = [5 \quad -5 \quad 0 \quad 1]$$

$$\text{linear}(X) = XW + b$$

$$\text{linear}\left(\begin{bmatrix} 1 & 2 & 1 \\ 3 & 4 & 0 \end{bmatrix}\right) \rightarrow \begin{bmatrix} 8 & 3 & 4 & 14 \\ 11 & 5 & 4 & 30 \end{bmatrix}$$

X shape=[batchsize, d_in]
W shape=[n, d_out]
b shape=[1, d_out]

MatMul

$$\begin{bmatrix} 1 & 2 & 1 \\ 3 & 4 & 0 \end{bmatrix} \begin{bmatrix} 2 & 2 & 0 & 3 \\ 0 & 1 & 1 & 5 \\ 1 & 4 & 2 & 0 \end{bmatrix} = \begin{bmatrix} 3 & 8 & 4 & 13 \\ 6 & 10 & 4 & 29 \end{bmatrix}$$

+ bias

$$\begin{bmatrix} 3 & 8 & 4 & 13 \\ 6 & 10 & 4 & 29 \end{bmatrix} + \begin{bmatrix} 5 & -5 & 0 & 1 \end{bmatrix} = \begin{bmatrix} 3 & 8 & 4 & 13 \\ 6 & 10 & 4 & 29 \end{bmatrix} + \begin{bmatrix} 5 & -5 & 0 & 1 \\ 5 & -5 & 0 & 1 \end{bmatrix} = \begin{bmatrix} 8 & 3 & 4 & 14 \\ 11 & 5 & 4 & 30 \end{bmatrix}$$

Broadcast sum

Bias broadcasting

Linear(X) takes an input with dimensionality `d_in` and transforms it to dimensionality `d_out`

$$\text{linear}(X) = XW + b$$

X shape=[batchsize, d_in]
W shape=[n, d_out]
b shape=[1, d_out]

+ bias $\begin{bmatrix} 3 & 8 & 4 & 13 \\ 6 & 10 & 4 & 29 \end{bmatrix} + \begin{bmatrix} 5 & -5 & 0 & 1 \end{bmatrix} = \begin{bmatrix} 3 & 8 & 4 & 13 \\ 6 & 10 & 4 & 29 \end{bmatrix} + \begin{bmatrix} 5 & -5 & 0 & 1 \\ 5 & -5 & 0 & 1 \end{bmatrix} = \begin{bmatrix} 8 & 3 & 4 & 14 \\ 11 & 5 & 4 & 30 \end{bmatrix}$

Broadcast sum

Tip: ndarray libraries like numpy and pytorch will not actually copy the bias vector like this. As an optimization, they implement broadcast semantics that achieves this behavior without having to explicitly copy the bias vector.

If you wanted to write
 $\text{linear}(X)$ more formally:

$$\text{linear}(X) = XW + \mathbf{1}b$$

Where $\mathbf{1}$ has shape=[batchsize, 1]

Math-y way of saying: "create 'batchsize' copies of b , row-wise"

Fully connected networks

A L -layer, fully connected network, a.k.a. multi-layer perceptron (MLP), with an explicit bias term, is defined by the iteration


$$Z_{i+1} = \sigma_i(\text{Linear}_i(X)), \quad i = 1, \dots, L$$

$$h_\theta(X) \equiv Z_{L+1}$$

$$Z_1 \equiv X$$


$$\text{Linear}_i(X) = XW_i + b_i$$

with model parameters $\theta = \{W_{1:L}, b_{1:L}\}$, and where $\sigma_i(x)$ is the activation, usually with $\sigma_L(x) = x$ (eg no activation fn for final output)




Deep learning: modularity

Amazingly, the Linear and CrossEntropyLoss layers are ubiquitous layers in deep learning, including: image/text classification, large language models, etc.



Resnet paper, ["Deep Residual Learning for Image Recognition"](#).



Transformers paper: <https://arxiv.org/pdf/1706.03762>

Figure 1: The Transformer - model architecture.

Deep learning: aka building models with modular "lego blocks".

Common layers: Linear, CrossEntropy, EWiseAdd, Conv2D, ...)

Key questions for fully connected networks

In order to actually train a fully-connected network (or any deep network), we need to address a certain number of questions:

- How do we choose the width and depth of the network?
- How do we actually optimize the objective? (“SGD” is the easy answer, but not the algorithm most commonly used in practice)
- How do we initialize the weights of the network?
- How do we ensure the network can continue to be trained easily over multiple optimization iterations?

All related questions that affect each other

There are (still) no definite answers to these questions, and for deep learning they wind up being problem-specific, but we will cover some basic principles

Outline

Fully connected networks

Optimization

Initialization

Optimization

Definition: an **objective function** is a scalar-valued function $f: \mathbb{R}^n \rightarrow \mathbb{R}$ that we would like to minimize.

Definition: an optimization problem consists of an objective function and optional constraints. We express it as:

$$\min_{\theta} f(\theta)$$

"Find the parameters θ that minimize $f(\theta)$ (aka "argmin")

Subject to: $g_i(\theta) \leq 0, i = 1, \dots, m$

$h_j(\theta) = 0, j = 1, \dots, p$



Fortunately, in this class (and deep learning in general) we typically don't have constraints, so we can drop these

Example: let $f(\theta, \mathcal{D}_{train}) = loss_{ce}(h(\theta), \mathcal{D}_{train})$.

Minimizing $f(\theta, \mathcal{D}_{train})$ is precisely what we do when we train an ML model!


Objective function: difficulty

What kinds of objective functions are "easy" to optimize? What makes a function "hard" to optimize?


$$\min_{\theta} f(\theta)$$

Subject to: $g_i(\theta) \leq 0, i = 1, \dots, m$


$h_j(\theta) = 0, j = 1, \dots, p$




"Easy"



"Hard"



Unbounded



Infeasible

$$f(\theta) \leq k$$

Global vs local optima



If you find a minima (say, through gradient descent), is it a global minima ("best")?

Or is it a local minima?

$$\min_{\theta} f(\theta)$$


Subject to: $g_i(\theta) \leq 0, i = 1, \dots, m$

$$h_j(\theta) = 0, j = 1, \dots, p$$



Convex functions


The optimization field has identified several "classes" of functions that are easy to optimize. One popular class of functions is called "convex functions".



By Varagk - Own work, CC BY-SA 4.0,
<https://commons.wikimedia.org/w/index.php?curid=124668501>



One lovely property of convex objective functions: if a convex function has a local optima, it is also a global optima*!



By Indeed123 - commons This diagram was created with Mathematica., CC BY-SA 3.0,
<https://commons.wikimedia.org/w/index.php?curid=5508870>

Gradient descent on convex objective functions will lead you to a global optima**!

Objective functions in deep learning

Unfortunately, the objective functions used in deep learning models are (very much) not convex.

Thus, many of the insights/techniques from convex optimization simply don't directly translate over.

Deep learning objective functions have many (many) local minima. The loss landscape is hard (impossible) to interpret, with very little guarantees.

Ex: gradient descent is only guaranteed to give you a **local** minima. No **global** minima guarantees.

Yet! Deep learning is still wildly successful despite giving up global optima and accepting local optima!




Figure 1: The loss surfaces of ResNet-56 with/without skip connections. The proposed filter normalization scheme is used to enable comparisons of sharpness/flatness between the two figures.

"Visualizing the Loss Landscape of Neural Nets", <https://arxiv.org/abs/1712.09913>

Optimization: closing thoughts

Optimization is a rich field with tons of interesting theory and applications.

$$\min_{\theta} f(\theta)$$

Subject to: $g_i(\theta) \leq 0, i = 1, \dots, m$

$$h_j(\theta) = 0, j = 1, \dots, p$$

In this course: we won't really study optimization that deeply. Mostly we'll study and apply gradient descent to train our models.

Gradient descent



In optimization, gradient descent is a popular and successful way to optimize functions. For a function f and iterate number t , the gradient descent update equation is:

$$\theta_{t+1} = \theta_t - \alpha \nabla_{\theta} f(\theta_t)$$

where $\alpha > 0$ is step size (learning rate), $\nabla_{\theta} f(\theta_t)$ is gradient evaluated at the parameters θ_t

Intuition: takes the “steepest descent direction” locally


Gradient $\nabla_{\theta} f(\theta_t)$ only tells you the **direction** to travel. Step size α dictates how "**far**" to go in that direction.



Contour plots


In optimization, contour plots are a way to visualize the loss landscape with respect to the optimization parameters.

Suppose our loss function is $y = f(x_1, x_2)$, where x_1, x_2, y are scalars



3D Plot

X1: param 1
X2: param 2
Y: loss value



Contour Plot


This circle (eg contour, or "level set") means:
"these values of x_1, x_2 lead to a loss value of 6"

Global optimum



Illustration of gradient descent

For $\theta \in \mathbb{R}^2$, consider quadratic function $f(\theta) = \frac{1}{2}\theta^T P\theta + q^T \theta$, for P positive definite (all positive eigenvalues)


Illustration of gradient descent with different step sizes:



Updates are a little erratic,
but still reaches the minimum




Smoother update trajectory.
A "better" learning rate



Gradient descent: local linear approximation

- **Gradient descent assumption:** approximate objective function via a locally-linear function.
 - Gradient $\nabla_{\theta} f(\theta_t)$ is tangent line of f at location θ_t . Tells us the direction that objective function decreases (within local linear window)
- **Pro:** fast, easy to compute!
- **Con:** locally-linear assumption can be (wildly!) inaccurate
- Sensitive to step size
 - Large step size \rightarrow "trust" local linear approx.
 - Smaller step size \rightarrow don't trust it too much



First order, second order methods

- Gradient descent is known as a "first order method", because it only uses local gradient information (aka "first order derivatives of the objective function")
- There exist a class of more sophisticated methods that utilize additional information (local curvature) to try to create a better (more accurate) local approximation of the loss function
 - These are called "second order methods"
 - The "additional curvature information" is: second-derivatives!

The Hessian

Definition: for a function $f: \mathbb{R}^n \rightarrow \mathbb{R}$ (eg a scalar-valued function), the **Hessian matrix** H (sometimes denoted ∇^2) is the $n \times n$ matrix of all second derivatives.

Intuition: Hessian describes local curvature of the function.

Vs **Jacobian** (aka the "gradient", as f is scalar-valued) gives you a linear approximation of the function at a point x (eg tangent line/plane/hyperplane for $\mathbb{R}^2, \mathbb{R}^3, \mathbb{R}^{n>3}$ respectively).

$$H = \nabla^2 f = \begin{bmatrix} \frac{\partial^2 f}{\partial x_1^2} & \frac{\partial^2 f}{\partial x_1 \partial x_2} & \dots & \frac{\partial^2 f}{\partial x_1 \partial x_n} \\ \frac{\partial^2 f}{\partial x_2 \partial x_1} & \frac{\partial^2 f}{\partial x_2^2} & \dots & \frac{\partial^2 f}{\partial x_2 \partial x_n} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial^2 f}{\partial x_n \partial x_1} & \frac{\partial^2 f}{\partial x_n \partial x_2} & \dots & \frac{\partial^2 f}{\partial x_n^2} \end{bmatrix}$$

Shape=[n, n]

$$J = \nabla f(x) = \begin{bmatrix} \frac{\partial f(x)}{\partial x_1} & \frac{\partial f(x)}{\partial x_2} & \dots & \frac{\partial f(x)}{\partial x_n} \end{bmatrix}$$

Shape=[1, n]

Newton's Method


One way to integrate more “global” structure into optimization methods is Newton’s method, which scales gradient according to inverse of the Hessian (matrix of second derivatives)

$$\theta_{t+1} = \theta_t - \alpha \left(\nabla_{\theta}^2 f(\theta_t) \right)^{-1} \nabla_{\theta} f(\theta_t)$$

where $\nabla_{\theta}^2 f(\theta_t)$ is the *Hessian*, $n \times n$ matrix of all second derivatives

Equivalent to approximating the function as quadratic using second-order Taylor expansion, then solving for optimal solution

Full step given by $\alpha = 1$, otherwise called a *damped* Newton method




A comparison of gradient descent (green) and Newton's method (red) for minimizing a function (with small step sizes). Newton's method uses curvature information (i.e. the second derivative) to take a more direct route.

By Oleg Alexandrov - self-made with en:Matlab. Tweaked in en:inkscape, Public Domain, <https://commons.wikimedia.org/w/index.php?curid=2284243>


(optional) to learn more, see:

https://en.wikipedia.org/wiki/Newton%27s_method_in_optimization

Linear vs Quadratic local approximations



Linear approximation
(in 3D, a 2D surface
or tangent plane)



Quadratic
approximation

Here is a neat visualization of linear vs quadratic local approximations: "[What do quadratic approximations look like?](#)"

(optional) If you're interested in learning more about linear/quadratic approximations and its application to optimization (eg Newton's method), see this "[Quadratic Approximation](#)" Khan Academy lesson.
(this is out-of-scope for this class though)

Illustration of Newton's method

Newton's method (with $\alpha = 1$) will optimize quadratic functions in one step

Not of that much practical relevance to deep learning for two reasons


1. We can't efficiently solve for Newton step, even using automatic differentiation (though there are tricks to approximately solve it)
2. For non-convex optimization, it's very unclear that we even *want* to use the Newton direction

In 2026: models with **hundreds of millions of parameters** are easily accessible.

Billion parameter models are easily productionized at scale at companies. (Ex: [Llama3-8B](#) has 8B params)

Trillion parameter models exist (albeit expensive to serve).

For a 1B param model: storing the Hessian matrix (1B x 1B) requires **4,000,000,000 gigabytes**. Not practical*!



*Disclaimer: I'm not fully familiar with modern second order methods. There exist sophisticated second order methods that don't have to explicitly construct the Hessian, but these approaches are outside the scope for this course

Back to first order methods



In deep learning (and this course!), we'll primarily stick to first order methods like gradient descent, with the following update rule:

$$\theta_{t+1} = \theta_t - \alpha \nabla_{\theta} f(\theta_t)$$

The above equation is known as "vanilla" gradient descent. Next, let's explore some common improvements to vanilla gradient descent!

Gradient descent deficiencies

- One problem with "vanilla" gradient descent:
noisy / oscillating gradients
 - **Possible root cause:** batchsize too small,
dataset not representative enough
- Another issue: getting stuck in local minima
- **Idea:** maintain a history of gradients. Perform
gradient update using a weighted average of
current gradient $\nabla_{\theta} f(\theta_t)$ and previous
gradients $[\nabla_{\theta} f(\theta_{t-1}), \nabla_{\theta} f(\theta_{t-2}), \dots]$



<https://optimization.cbe.cornell.edu/index.php?title=Momentum>

Momentum

This is known as the *momentum* update, that takes into account a moving average of *multiple* previous gradients

$$\begin{aligned} u_{t+1} &= \beta u_t + (1 - \beta) \nabla_{\theta} f(\theta_t) \\ \theta_{t+1} &= \theta_t - \alpha u_{t+1} \end{aligned}$$

where $\alpha \in \mathbb{R}$ is step size as before, and $\beta \in \mathbb{R}$ is momentum averaging parameter

- Note: often written in alternative forms $u_{t+1} = \beta u_t + \nabla_{\theta} f(\theta_t)$ (or $u_{t+1} = \beta u_t + \alpha \nabla_{\theta} f(\theta_t)$) but I prefer above to keep u the same “scale” as gradient

Momentum: gradient history

Momentum update:

$$\begin{aligned} u_{t+1} &= \beta u_t + (1 - \beta) \nabla_{\theta} f(\theta_t) \\ \theta_{t+1} &= \theta_t - \alpha u_{t+1} \end{aligned}$$

where $\alpha \in \mathbb{R}$ is step size as before, and $\beta \in \mathbb{R}$ is momentum averaging parameter

$$u_0 = 0$$

$$u_1 = \beta u_0 + (1 - \beta) \nabla_{\theta} f(\theta_0) = (1 - \beta) \nabla_{\theta} f(\theta_0)$$

$$u_2 = \beta u_1 + (1 - \beta) \nabla_{\theta} f(\theta_1) = \beta((1 - \beta) \nabla_{\theta} f(\theta_0)) + (1 - \beta) \nabla_{\theta} f(\theta_1)$$

$$\begin{aligned} u_3 &= \beta u_2 + (1 - \beta) \nabla_{\theta} f(\theta_2) = \beta(\beta((1 - \beta) \nabla_{\theta} f(\theta_0)) + (1 - \beta) \nabla_{\theta} f(\theta_1)) + (1 - \beta) \nabla_{\theta} f(\theta_2) \\ &= \beta^2(1 - \beta) \nabla_{\theta} f(\theta_0) + \beta(1 - \beta) \nabla_{\theta} f(\theta_1) + (1 - \beta) \nabla_{\theta} f(\theta_2) \end{aligned}$$

...

$$u_n = \underbrace{\beta^{n-1}(1 - \beta) \nabla_{\theta} f(\theta_0) + \beta^{n-2}(1 - \beta) \nabla_{\theta} f(\theta_1) + \cdots + (1 - \beta) \nabla_{\theta} f(\theta_{n-1})}_{\text{Exponential decay. Exponential "forgetting" of previous gradients, where memory strength is controlled by } \beta} + \underbrace{\nabla_{\theta} f(\theta_n)}_{\text{Current gradient}}$$

β is typically between 0.8 and 0.99. 0.9 is a decent starting value.





Exponential decay. Exponential "forgetting" of previous gradients, where memory strength is controlled by β .

Current gradient

Illustration of momentum

Momentum “smooths” out the descent steps, but can also introduce other forms of oscillation and non-descent behavior

Frequently useful in training deep networks in practice



“Unbiasing” momentum terms

The momentum term u_t (if initialized to zero, as is common), will have smaller magnitude in initial iterations than in later ones. Results in the first few iterations taking smaller steps than vanilla GD, resulting in slower convergence.





To “unbias” the update to have equal expected magnitude across all iterations, we can use the update

$$\begin{aligned} u_{t+1} &= \beta u_t + (1 - \beta) \nabla_{\theta} f(\theta_t) \\ \theta_{t+1} &= \theta_t - \alpha u_{t+1} \end{aligned}$$



$$\begin{aligned} u_{t+1} &= \beta u_t + (1 - \beta) \nabla_{\theta} f(\theta_t) \\ \theta_{t+1} &= \theta_t - \alpha \frac{u_{t+1}}{(1 - \beta^{t+1})} \end{aligned}$$

Note: here, β^{t+1} means β raised to the power of $(t+1)$







Nesterov Momentum

One (admittedly, of many) useful tricks is the notion of Nesterov momentum (or Nesterov acceleration), which computes momentum update at “next” point

$$\begin{aligned} u_{t+1} &= \beta u_t + (1 - \beta) \nabla_{\theta} f(\theta_t) \\ \theta_{t+1} &= \theta_t - \alpha u_{t+1} \end{aligned} \Rightarrow \begin{aligned} u_{t+1} &= \beta u_t + (1 - \beta) \nabla_{\theta} f(\theta_t - \alpha u_t) \\ \theta_{t+1} &= \theta_t - \alpha u_{t+1} \end{aligned}$$

"look ahead" one step

A “good” thing for convex optimization, and (sometimes) helps for deep networks



Feels like a simple tweak, but for convex functions this provably speeds up convergence ($1/k$ to $1/k^2$, where k is num iterations). Neat!

Adam

The *scale* of the gradients can vary widely for different parameters, especially e.g. across different layers of a deep network, different layer types, etc

So-called *adaptive gradient* methods attempt to estimate this scale over iterations and then re-scale the gradient update accordingly

Most widely used adaptive gradient method for deep learning is Adam algorithm, which combines momentum and adaptive scale estimation

$$\begin{aligned} u_{t+1} &= \beta_1 u_t + (1 - \beta_1) \nabla_{\theta} f(\theta_t) \\ v_{t+1} &= \beta_2 v_t + (1 - \beta_2) (\nabla_{\theta} f(\theta_t))^2 \\ \theta_{t+1} &= \theta_t - \alpha \frac{u_{t+1}}{\left(\sqrt{v_{t+1}} + \epsilon \right)} \end{aligned}$$

Square is done elementwise

Division is done elementwise

Square root is done elementwise

Adam with bias correction

Common to apply bias correction as well:

$$\begin{aligned} u_{t+1} &= \beta_1 u_t + (1 - \beta_1) \nabla_{\theta} f(\theta_t) \\ v_{t+1} &= \beta_2 v_t + (1 - \beta_2) (\nabla_{\theta} f(\theta_t))^2 \\ \theta_{t+1} &= \theta_t - \alpha \frac{u_{t+1}}{(\sqrt{v_{t+1}} + \epsilon)} \end{aligned}$$



$$\begin{aligned} u_{t+1} &= \beta_1 u_t + (1 - \beta_1) \nabla_{\theta} f(\theta_t) \\ \hat{u}_{t+1} &= \frac{u_{t+1}}{(1 - \beta_1^{t+1})} \\ v_{t+1} &= \beta_2 v_t + (1 - \beta_2) (\nabla_{\theta} f(\theta_t))^2 \\ \hat{v}_{t+1} &= \frac{v_{t+1}}{(1 - \beta_2^{t+1})} \\ \theta_{t+1} &= \theta_t - \alpha \frac{\hat{u}_{t+1}}{(\sqrt{\hat{v}_{t+1}} + \epsilon)} \end{aligned}$$

Typical choices: $\beta_1 = 0.9$, $\beta_2 = 0.999$.


Means: variance estimate \hat{v}_{t+1} moves much more slowly than the momentum term \hat{u}_{t+1} .

$$\epsilon = 10^{-6}.$$

Notes on / illustration of Adam

Whether Adam is “good” optimizer is endlessly debated within deep learning, but it often seems to work quite well in practice (maybe?)

There are alternative universes where endless other variants became the “standard” (no unbiasing? average of absolute magnitude rather than squared? Nesterov-like acceleration?) but Adam is well-tuned and hard to uniformly beat



"Stochastic" gradient descent

"Standard" gradient descent would run through your entire dataset to calculate the gradient.

However, in deep learning, it's common to instead calculate your gradient over a subset of the dataset, called a "minibatch". This is known as "stochastic" gradient descent.

Fortunately, the math for "full dataset" vs "minibatch" gradient descent works out: both produce unbiased estimates of the gradient. For practical reasons, we do "minibatch" (aka "stochastic") gradient descent.

	Quality of gradient $\nabla_{\theta} f(\theta_t)$	Speed
"Full batch" gradient descent	Better	Slower
"Stochastic" gradient descent	Worse (noisier)	Faster
"Stochastic" gradient descent with tiny batchsize (say, batchsize=1)	Worst (noisiest)	Fastest





SGD: quality vs speed tradeoff

This leads us again to the SGD algorithm, repeating for batches $B \subset \{1, \dots, m\}$

$$\theta_{t+1} = \theta_t - \frac{\alpha}{B} \sum_{i \in B} \nabla_{\theta} \ell(h(x^{(i)}), y^{(i)})$$

Tip: since decreasing batchsize leads to noisier gradients, we typically reduce learning rate accordingly (trust gradient less \rightarrow take smaller step)

Instead of taking a few expensive, noise-free, steps, we take *many* cheap, noisy steps, which ends having much strong performance per compute



The most important takeaways

All the optimization methods you have seen thus far presented are *only* actually used in their stochastic form

The amount of valid intuition about these optimization methods you will get from looking at simple (convex, quadratic) optimization problems is limited

You need to constantly experiment to gain an understanding / intuition of how these methods actually affect deep networks of different types

Outline

Fully connected networks

Optimization

Initialization

Initialization of weights

Recall that we optimize parameters iteratively by stochastic gradient descent, e.g.

$$W_i := W_i - \alpha \nabla_{W_i} \ell(h_\theta(X), y)$$

But how do we choose the *initial* values of W_i , b_i ? (maybe just initialize to zero?)

Recall the manual backpropagation forward/backward passes (without bias):

$$\begin{aligned} Z_{i+1} &= \sigma_i(Z_i W_i) \\ G_i &= (G_{i+1} \circ \sigma'_i(Z_i W_i)) W_i^T \end{aligned}$$

- If $W_i = 0$, then $G_j = 0$ for $j \leq i$, $\Rightarrow \nabla_{W_i} \ell(h_\theta(X), y) = 0$
- I.e., $W_i = 0$ is a (very bad) local optimum of the objective


Key idea #1: Choice of initialization matters

Let's just initialize weights "randomly", e.g., $W_i \sim \mathcal{N}(0, \sigma^2 I)$

The choice of variance σ^2 will affect two (related) quantities:

1. The norm of the forward activations Z_i
2. The norm of the gradients $\nabla_{W_i} \ell(h_\theta(X), y)$

Illustration on MNIST with
 $n = 100$ hidden units,
depth 50, ReLU
nonlinearities



$\sigma^2 = \frac{1}{n}$: activations shrink with model depth
(bad! Can lead to slow training, "vanishing gradient problem")

$\sigma^2 = \frac{2}{n}$ is the "best" here:
stable activation norms, and
"healthy" gradient norm.

$\sigma^2 = \frac{3}{n}$: activations grow with model depth (bad!
Can lead to exploding gradients, which can lead to overshooting and training instability)

Key idea #2: Weights don't move “that much”

Might have the picture in your mind that the parameters of a network converge to some similar region of points regardless of their initialization

This is not true ... weights often stay much closer to their initialization than to the “final” point after optimization from different

End result: initialization matters ... we'll see some of the practical aspects next lecture

Linear layer initialization strategy (1/2)

Suppose we have a linear layer L parameterized with weight matrix W , defined as:
 $L(x) = xW$, where x has shape=[1, n], and W has shape=[n, m].

Suppose the inputs to L are normally distributed: $x \sim \mathcal{N}(0,1)$

Goal: to keep activations "healthy" (eg don't grow/shrink with model depth), we'd like the outputs $L(x)$ to also be normally distributed: $L(x) \sim \mathcal{N}(0,1)$

Question: how do we initialize W to achieve $L(x) \sim \mathcal{N}(0,1)$?

Hint: let $w \sim \mathcal{N}(0, \sigma^2)$. What would be a "good" value for σ^2 ?

Linear layer initialization strategy (2/2)

Question: how do we initialize W to achieve $L(x) \sim \mathcal{N}(0,1)$?

Hint: let $w \sim \mathcal{N}(0, \sigma^2)$. What would be a "good" value for σ^2 ?

Answer: $\sigma^2 = \frac{1}{n}$

n is shape of x

Proof: Consider independent random variables $x \sim \mathcal{N}(0,1)$, $w \sim \mathcal{N}\left(0, \frac{1}{n}\right)$; then

$$\mathbf{E}[x_i w_i] = \mathbf{E}[x_i] \mathbf{E}[w_i] = 0, \quad \mathbf{Var}[x_i w_i] = \mathbf{Var}[x_i] \mathbf{Var}[w_i] = 1/n$$

Since x, w are independent

so $\mathbf{E}[w^T x] = 0$, $\mathbf{Var}[w^T x] = 1$ ($w^T x \rightarrow \mathcal{N}(0,1)$ by central limit theorem)

Thus, informally speaking, $L(x)$ achieves our goal: $x_i \sim \mathcal{N}(0, I)$, $W_i \sim \mathcal{N}\left(0, \frac{1}{n}I\right)$ then $L(x) = W_i^T z_i \sim \mathcal{N}(0, I)$

If we use a ReLU nonlinearity: then "half" the components of $L(x)$ will be set to zero, so we need double the variance on W_i to achieve the same final variance, hence $W_i \sim \mathcal{N}\left(0, \frac{2}{n}I\right)$ ([Kaiming normal initialization](#))